

Query Processing

Given a query,

```
SELECT B.manufacturer
FROM Beers B, Sells S
WHERE B.name = S.beer
AND S.price < 20
```

We execute it in two levels:

- ① Logical Plan : Relational Alg., many variations exist
- ② Physical Plan : Implementation of Logical Plan

Implementations of Operations

I. SELECTION

- ① Logical level: subject = "de" (Textbook)
- ② Physical Level
 - Table Scan : Read R block by block
 - Index Scan : Use Index to find blocks (efficient)

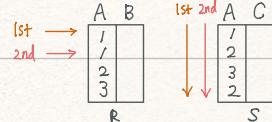
2. JOIN ↗ Can be extended to implement other operators

① Logical Level: $R \bowtie S$

② Physical Level :

CASE 1. Nested Loop Join

- ↳ If both R and S can fit into the main memory
- ↳ Time: $O(|R||S|)$



Basically, for i in $R(A)$
for j in $S(A)$

↳ Memory Requirement: ④

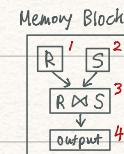
↳ Cost: $B(R) + |R|B(S)$

Read blockwise

$$B(R) + B(R) \cdot B(S) \approx \frac{B(R)B(S)}{M}$$

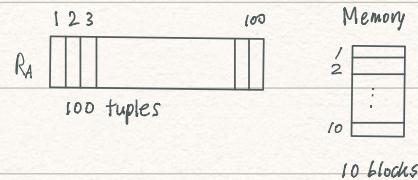
Improved w/ bigger M

If $M = 200$ (large),
You can fill 199 blocks w/ R
and join w/ S



Internal Memory Join

* Problem: how to sort R if it's too big to fit in memory?

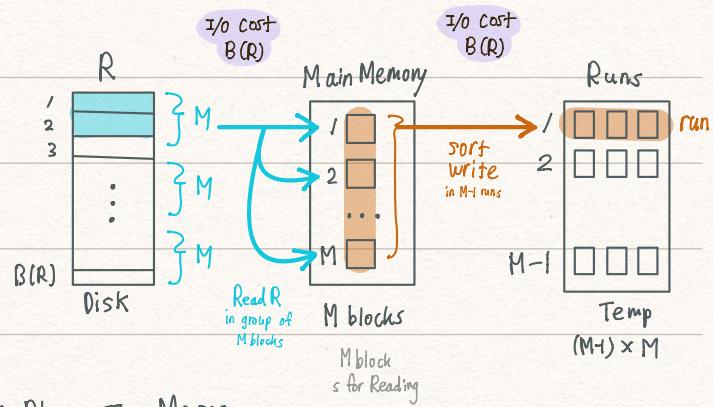


Case 3. Two Pass (Phase) Multiway Merge Sort

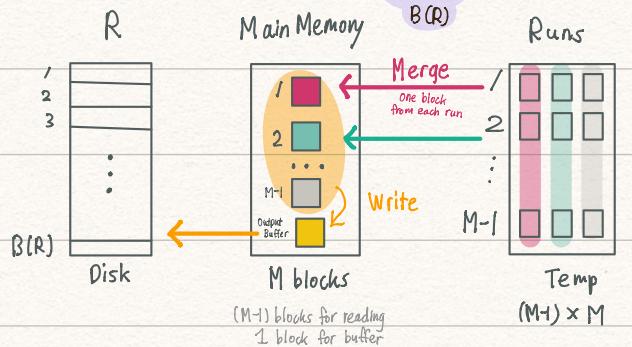
↳ At least one relation doesn't fit into main memory.

↳ Dominant cost is I/O access

• Phase I: Read & Sort



• Phase II: Merge

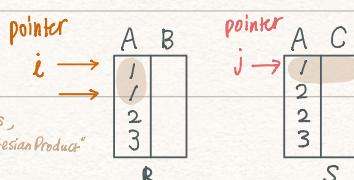


↳ Cost: $2B(R)$ in Phase I $\rightarrow 3B(R)$
+ $B(R)$ in Phase II

↳ Limitation: $B(R) \leq M(M-1)$
or $B(R) \leq M^2$

CASE 2. Sort-Merge Join

- ↳ If both R and S can fit into the main memory
- ↳ $O(|R| \log |R|)$ if $|R| > |S|$



for duplicates,
we use "Cartesian Product"

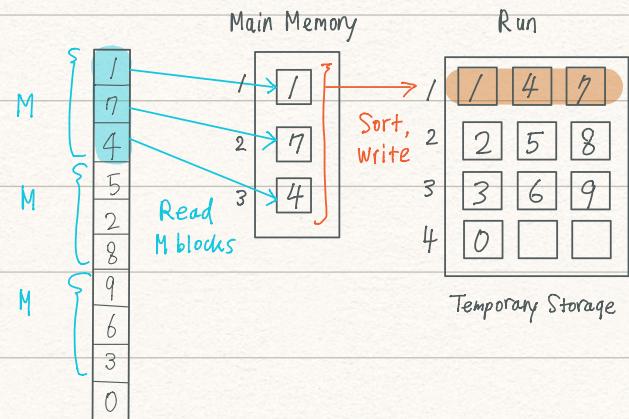
Since $R(A)$ and $S(A)$ are sorted,
we can search and match.

External Memory Join

Case 4. General Multipass Multiway Merge Sort

ex> M : 3 blocks, each block holding one number

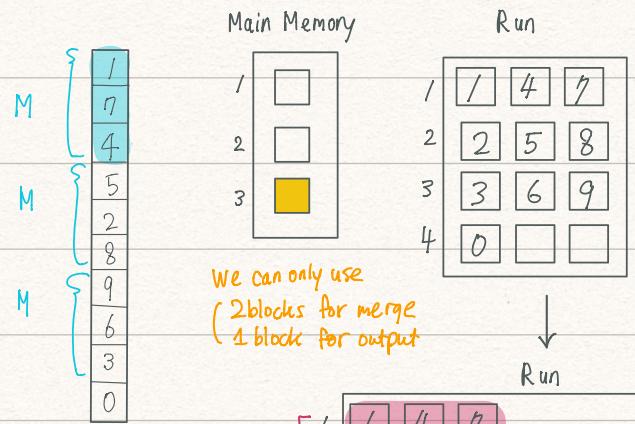
- Phase I (Pass 0) : Read & Sort



- Phase II (Pass 1) : Merge

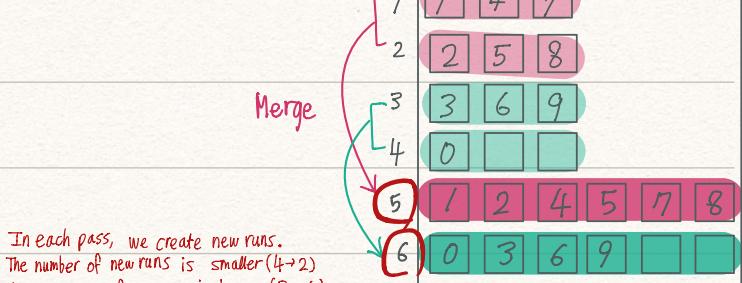
The key to apply Multipass Multiway Merge Sort is
 ↗ Can handle more than $B(R) > M^2$

- For each pass, merge as much as you can w/ given M
- Create new runs (Merge & Sort) w/ each run # runs ↓, |run| ↑
- Finish ② when everything is sorted.
- Return the one sorted relation.

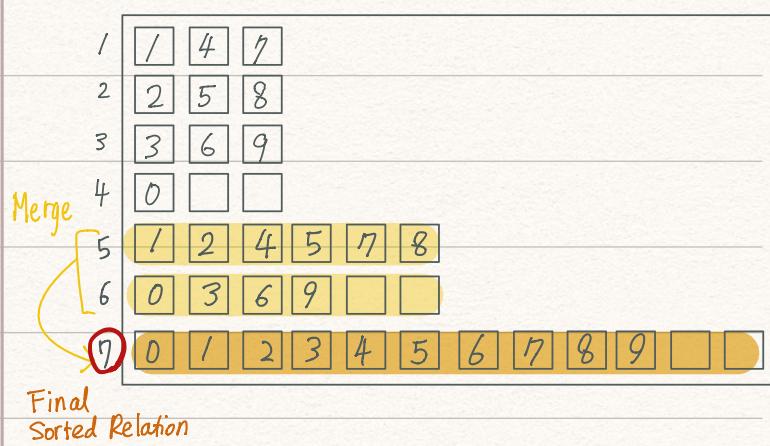


Generalization of MMMS :

- Pass 0 : Read a relation, R, in group of M blocks & sort them
 $\hookrightarrow \lceil \frac{B(R)}{M} \rceil$ runs in level 0.
- Pass k : Merge $(M-1)$ runs in level $(k-1)$ & write them in level k
 why? $(M-1)$ blocks for input
 1 block for output
 \hookrightarrow In level k, the number of runs = The number of runs in level $(k-1)$ $\frac{(M-1)}{(M-1)}$
- Final Pass : results in one sorted run.



- Phase III (Pass 2) : Merge



Analysis of MMMS :

- # of Passes : $\lceil \log_{M-1} \lceil \frac{B(R)}{M} \rceil \rceil + 1$
Pass 0 for reading R
- Cost - In each pass, we read and write R
 \hookrightarrow # of passes • $2B(R)$
 - In the final pass, we simply write, so disregard
 \therefore cost of MMMS = $O(B(R) \cdot \log_M B(R))$
- Memory Requirement: M