

# CSE 544: Principles of Database Systems

Conjunctive Queries  
(and beyond)

# Outline

- Conjunctive queries
- Query containment and equivalence
- Query minimization
- Undecidability for relational queries

# Conjunctive Queries (CQ)

- CQ = one datalog rule
- CQ = SELECT-DISTINCT-FROM-WHERE
- CQ = select/project/join ( $\sigma$ ,  $\Pi$ ,  $\bowtie$ ) fragment of RA
- CQ = existential/conjunctive ( $\exists$ ,  $\wedge$ ) fragment of RC

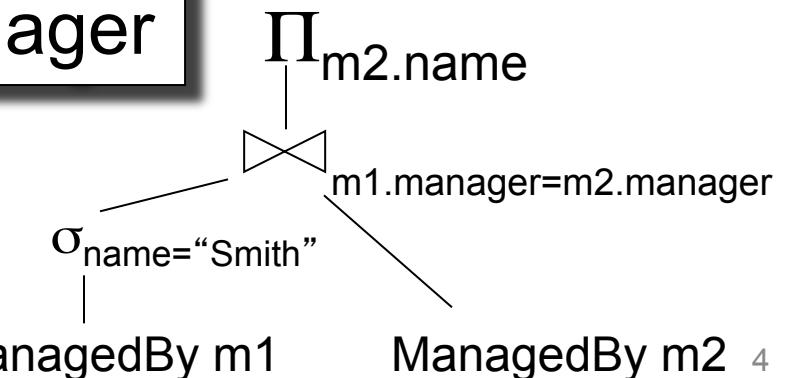
Notice: strictly speaking we are not allowed to use  $<$ ,  $\leq$ ,  $>$ ,  $\geq$ ,  $\neq$  in CQ's.  
If we include these, then the language is called CQ $^<$ , or CQ $^\neq$ , etc.

# Examples

Find all employees with same manager as “Smith”:

```
q(x) :- ManagedBy("Smith",y), ManagedBy(x,y)
```

```
SELECT DISTINCT m2.name  
FROM ManagedBy m1,  
      ManagedBy m2  
WHERE m1.name="Smith"  
  AND m1.manager=m2.manager
```



# Examples

- Example of CQ

$$q(x,y) = \exists z.(R(x,z) \wedge \exists u.(R(z,u) \wedge R(u,y)))$$

$$q(x) = \exists z.\exists u.(R(x,z) \wedge R(z,u) \wedge R(u,y))$$

- Examples of non-CQ:

$$q(x,y) = \forall z.(R(x,z) \rightarrow R(y,z))$$

$$q(x) = T(x) \vee \exists z.S(x,z)$$

# Query Equivalence and Containment

Containment/equivalence are examples of **static analysis**

Used in:

- Query optimization
- Query rewriting using views

# Query Equivalence

**Definition.** Queries  $q_1$  and  $q_2$  are **equivalent** if for every database  $D$ ,  $q_1(D) = q_2(D)$ .

Notation:  $q_1 \equiv q_2$

# Query Containment

**Definition.** Query  $q_1$  is **contained** in  $q_2$  if for every database  $D$ ,  $q_1(D) \subseteq q_2(D)$ .

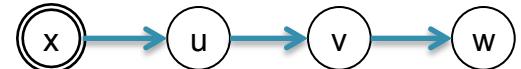
Notation:  $q_1 \subseteq q_2$

**Fact:**  $q_1 \subseteq q_2$  and  $q_2 \subseteq q_1$  iff  $q_1 = q_2$

We discuss only query containment.

# Example 1

Is  $q_1 \subseteq q_2$  ?

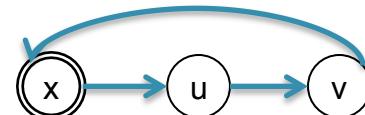


$$q_1(x) :- R(x,u), R(u,v), R(v,w)$$
$$q_2(x) :- R(x,u), R(u,v)$$

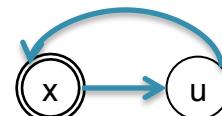


# Example 2

Is  $q_1 \subseteq q_2$  ?



```
q1(x) :- R(x,u), R(u,v), R(v,x)  
q2(x) :- R(x,u), R(u,x)
```



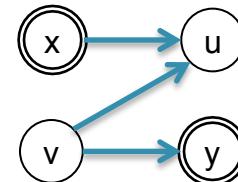
# Example 3

Is  $q_1 \subseteq q_2$  ?



$q_1(x,y) :- R(x,y)$

$q_2(x,y) :- R(x,u), R(v,u), R(v,y)$



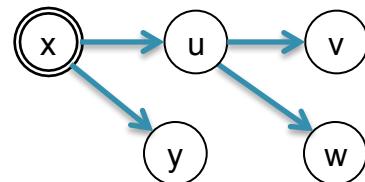
# Example 4

Is  $q_1 \subseteq q_2$  ?



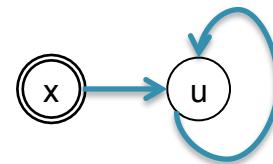
$q_1(x) :- R(x,u), R(u,v)$

$q_2(x) :- R(x,u), R(x,y), R(u,v), R(u,w)$

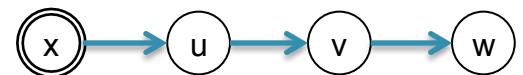


# Example 5

Is  $q_1 \subseteq q_2$  ?



```
q1(x) :- R(x,u), R(u,u)  
q2(x) :- R(x,u), R(u,v), R(v,w)
```



# Example 6

Is  $q_1 \subseteq q_2$  ?

```
q1(x) :- R(x,u), R(u,"Smith")
```

```
q2(x) :- R(x,u), R(u,v)
```

# Discussion

- We discuss only query containment for Boolean queries
  - If  $q_1, q_2$  are Boolean queries, then containment  $q_1(D) \subseteq q_2(D)$  means implication  $q_1(D) \rightarrow q_2(D)$
- If  $q_1, q_2$  are not Boolean, then convert them into Boolean queries by pretending that their head variables are constants
  - We must make *the same* constants in both  $q_1, q_2$

# Query Containment for CQ

**Theorem** [Chandra&Harel'1977]

Given two CQ,  $q_1$ ,  $q_2$ , the following are equivalent:

1. For every database  $D$ ,  $q_1(D) \subseteq q_2(D)$
2. There exists a homomorphism  $h : q_2 \rightarrow q_1$
3.  $q_2$  is true on the canonical database of  $q_1$

Moreover, this problem is NP-complete

# Query Homomorphisms

**Definition** A homomorphism  $h : q_2 \rightarrow q_1$  is a function  
 $h: \text{var}(q_2) \rightarrow \text{var}(q_1) \cup \text{const}(q_1)$   
such that, for every atom  $R(x, y, z, \dots)$  in the query  $q_2$ ,  
there is an atom  $R(h(x), h(y), h(z), \dots)$  in the query  $q_1$

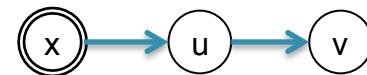
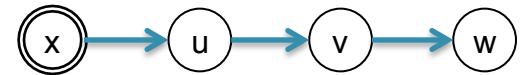
**The Homomorphism Theorem**  $q_1 \subseteq q_2$  iff there exists a homomorphism  $h : q_2 \rightarrow q_1$

# Example 1

Is  $q_1 \subseteq q_2$  ?

$q_1(x) :- R(x,u), R(u,v), R(v,w)$

$q_2(x) :- R(x,u), R(u,v)$

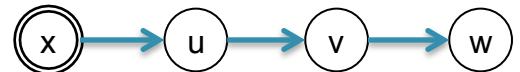


# Example 1

Is  $q_1 \subseteq q_2$  ?

$q_1(x) :- R(x,u), R(u,v), R(v,w)$

$q_2(x) :- R(x,u), R(u,v)$



First, transform the head variables into a constant: must be the same in  $q_1$  and  $q_2$ !

$q_1 :- R("x",u), R(u,v), R(v,w)$

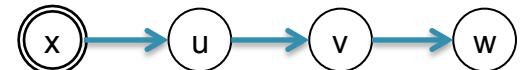
$q_2 :- R("x",u), R(u,v)$

# Example 1

Is  $q_1 \subseteq q_2$  ?

$q_1(x) :- R(x,u), R(u,v), R(v,w)$

$q_2(x) :- R(x,u), R(u,v)$



First, transform the head variables into a constant: must be the same in  $q_1$  and  $q_2$ !

$q_1 :- R("x",u), R(u,v), R(v,w)$

$q_2 :- R("x",u), R(u,v)$

The homomorphism  $h : q_2 \rightarrow q_1$  is the following:  
 $h(u) = u, h(v) = v$

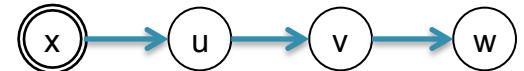
In class: check that  $h$  is a homomorphism

# Example 1

Is  $q_1 \subseteq q_2$  ?

$q_1(x) :- R(x,u), R(u,v), R(v,w)$

$q_2(x) :- R(x,u), R(u,v)$



First, transform the head variables into a constant: must be the same in  $q_1$  and  $q_2$ !

$q_1 :- R("x",u), R(u,v), R(v,w)$

$q_2 :- R("x",u), R(u,v)$

The homomorphism  $h : q_2 \rightarrow q_1$  is the following:  
 $h(u) = u, h(v) = v$

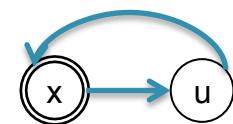
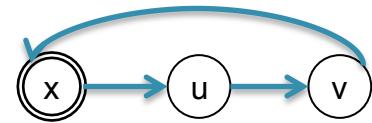
This proves  $q_1 \subseteq q_2$

In class: check that  $h$  is a homomorphism

# Example 2

Is  $q_1 \subseteq q_2$  ?

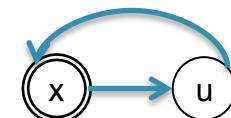
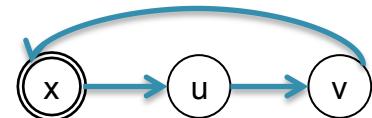
```
q1(x) :- R(x,u), R(u,v), R(v,x)  
q2(x) :- R(x,u), R(u,x)
```



# Example 2

Is  $q_1 \subseteq q_2$  ?

```
q1(x) :- R(x,u), R(u,v), R(v,x)  
q2(x) :- R(x,u), R(u,x)
```



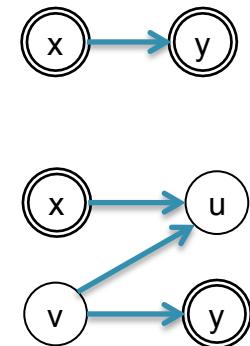
In class: check that there is no homomorphism

This proves  $q_1 \not\subseteq q_2$

# Example 3

Is  $q_1 \subseteq q_2$  ?

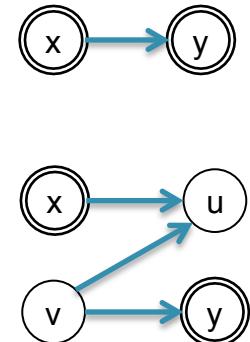
```
q1(x,y) :- R(x,y)  
q2(x,y) :- R(x,u), R(v,u), R(v,y)
```



# Example 3

Is  $q_1 \subseteq q_2$  ?

```
q1(x,y) :- R(x,y)  
q2(x,y) :- R(x,u), R(v,u), R(v,y)
```



Head variables x,y are treated like constants

The homomorphism  $h : q_2 \rightarrow q_1$  is the following:

$$h(u) = y, \quad h(v) = x$$

This proves  $q_1 \subseteq q_2$

# Canonical Database

**Definition.** Given a conjunctive query  $q$ ,  
the canonical database  $D_q$  is the following:

- The active domain =  $\text{vars}(q) \cup \text{const}(q)$
- Each atom in  $q$  becomes a tuple in  $D_q$

**The Canonical Database Theorem**  $q_1 \subseteq q_2$  iff  
the query  $q_2$  is true on the canonical database  $D_{q_1}$

# Example 1

Is  $q_1 \subseteq q_2$  ?

```
q1(x) :- R(x,u), R(u,v), R(v,w)  
q2(x) :- R(x,u), R(u,v)
```

# Example 1

Is  $q_1 \subseteq q_2$  ?

```
q1(x) :- R(x,u), R(u,v), R(v,w)  
q2(x) :- R(x,u), R(u,v)
```

$D_{q_1}$

R:

x	u
u	v
v	w

Compute  $q_2$  on  $D_{q_1}$ : the answer is **true**  
Keep in mind that “x” in  $q_2$  is a constant, it can only match x.

This proves  $q_1 \subseteq q_2$

# Example 2

Is  $q_1 \subseteq q_2$  ?

```
q1(x) :- R(x,u), R(u,v), R(v,x)  
q2(x) :- R(x,u), R(u,x)
```

# Example 2

Is  $q_1 \subseteq q_2$  ?

```
q1(x) :- R(x,u), R(u,v), R(v,x)  
q2(x) :- R(x,u), R(u,x)
```

$D_{q1}$

R:

x	u
u	v
v	x

Compute  $q_2$  on  $D_{q1}$ : the answer is **false**

This proves that  $q_1 \not\subseteq q_2$

# Query Containment for CQ

**Theorem** [Chandra&Harel'1977]

Given two CQ,  $q_1$ ,  $q_2$ , the following are equivalent:

1. For every database  $D$ ,  $q_1(D) \subseteq q_2(D)$
2. There exists a homomorphism  $h : q_2 \rightarrow q_1$
3.  $q_2$  is true on the canonical database of  $q_1$

Moreover, this problem is NP-complete

Proof: show in class that  $1 \rightarrow 3 \rightarrow 2 \rightarrow 1$

# The Complexity

**Theorem** Checking containment of two CQ queries is NP-complete

# Proof

Reduction from 3SAT

Given a 3CNF  $\Phi$

**Step 1:**

construct  $q_1$  independently of  $\Phi$ .

# Proof

Reduction from 3SAT

Given a 3CNF  $\Phi$

**Step 1:**

construct  $q_1$  independently of  $\Phi$ .

**Step 2:**

construct  $q_2$  from  $\Phi$ .

# Proof

## Reduction from 3SAT

Given a 3CNF  $\Phi$

**Step 1:**

construct  $q_1$  independently of  $\Phi$ .

**Step 2:**

construct  $q_2$  from  $\Phi$ .

**Claim:**

there exists a

homomorphism  $q_2 \rightarrow q_1$

iff  $\Phi$  is satisfiable

# Proof

Reduction from 3SAT

Given a 3CNF  $\Phi$

**Step 1:**

construct  $q_1$  independently of  $\Phi$ .

**Step 2:**

construct  $q_2$  from  $\Phi$ .

**Claim:**

there exists a  
homomorphism  $q_2 \rightarrow q_1$   
iff  $\Phi$  is satisfiable

Running example:

$$\begin{aligned}\Phi = & (\neg X_3 \vee \neg X_1 \vee X_4) \\ & (X_1 \vee X_2 \vee X_3) \\ & (\neg X_2 \vee \neg X_3 \vee X_1)\end{aligned}$$

# Step 1: Constructing $q_1$

There are four types  
of clauses possible  
in a 3SAT:

$$\text{Type 1} = \neg X \vee \neg Y \vee \neg Z$$

# Step 1: Constructing $q_1$

There are four types  
of clauses possible  
in a 3SAT:

$$\text{Type 1} = \neg X \vee \neg Y \vee \neg Z$$

$$\text{Type 2} = \neg X \vee \neg Y \vee Z$$

# Step 1: Constructing $q_1$

There are four types  
of clauses possible  
in a 3SAT:

$$\text{Type 1} = \neg X \vee \neg Y \vee \neg Z$$

$$\text{Type 2} = \neg X \vee \neg Y \vee Z$$

$$\text{Type 3} = \neg X \vee Y \vee Z$$

# Step 1: Constructing $q_1$

There are four types  
of clauses possible  
in a 3SAT:

$$\text{Type 1} = \neg X \vee \neg Y \vee \neg Z$$

$$\text{Type 2} = \neg X \vee \neg Y \vee Z$$

$$\text{Type 3} = \neg X \vee Y \vee Z$$

$$\text{Type 4} = X \vee Y \vee Z$$

For each type, we include in  $q_1$  a relation with all 7 satisfying assignments

R

(misses 1,1,1)

0	0	0
0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	0

# Step 1: Constructing $q_1$

There are four types  
of clauses possible  
in a 3SAT:

$$\text{Type 1} = \neg X \vee \neg Y \vee \neg Z$$

$$\text{Type 2} = \neg X \vee \neg Y \vee Z$$

$$\text{Type 3} = \neg X \vee Y \vee Z$$

$$\text{Type 4} = X \vee Y \vee Z$$

For each type, we include in  $q_1$  a relation with all 7 satisfying assignments

**R**  
(misses 1,1,1)

**S**  
(misses 1,1,0)

0	0	0
0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	0

0	0	0
0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	1

# Step 1: Constructing $q_1$

There are four types of clauses possible in a 3SAT:

$$\text{Type 1} = \neg X \vee \neg Y \vee \neg Z$$

$$\text{Type 2} = \neg X \vee \neg Y \vee Z$$

$$\text{Type 3} = \neg X \vee Y \vee Z$$

$$\text{Type 4} = X \vee Y \vee Z$$

For each type, we include in  $q_1$  a relation with all 7 satisfying assignments

**R**  
(misses 1,1,1)

**S**  
(misses 1,1,0)

**T**  
(misses 1,0,0)

**K**  
(misses 0,0,0)

0	0	0
0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	0

0	0	0
0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	1

0	0	0
0	0	1
0	1	0
0	1	1
1	0	1
1	1	0
1	1	1

0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	0
1	1	1

# Step 1: Constructing $q_1$

There are four types of clauses possible in a 3SAT:

$$\text{Type 1} = \neg X \vee \neg Y \vee \neg Z$$

$$\text{Type 2} = \neg X \vee \neg Y \vee Z$$

$$\text{Type 3} = \neg X \vee Y \vee Z$$

$$\text{Type 4} = X \vee Y \vee Z$$

For each type, we include in  $q_1$  a relation with all 7 satisfying assignments

R  
(misses 1,1,1)

S  
(misses 1,1,0)

T  
(misses 1,0,0)

K  
(misses 0,0,0)

0	0	0
0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	0

0	0	0
0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	1

0	0	0
0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	0

0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	0
1	1	1

So which are the atoms in  $q_1$ ?

# Step 2: Constructing $q_2$

Given  $\Phi$ , we construct  $q_2$  as follows.

For each clause in  $\Phi$  we create an atom in  $q_2$ :

- If the clause is of Type 1, then the atom is  $R$ ;  
if the clause is of Type 2, then the atom is  $S$ ; etc.
- The variables in this atom are the same as the variables in the clause

# Step 2: Constructing $q_2$

Given  $\Phi$ , we construct  $q_2$  as follows.

For each clause in  $\Phi$  we create an atom in  $q_2$ :

- If the clause is of Type 1, then the atom is  $R$ ;  
if the clause is of Type 2, then the atom is  $S$ ; etc.
- The variables in this atom are the same as the variables in the clause

---

Example:

$$\Phi = (\neg X_3 \vee \neg X_1 \vee X_4) \wedge (X_1 \vee X_2 \vee X_3) \wedge (\neg X_2 \vee \neg X_3 \vee X_1)$$

# Step 2: Constructing $q_2$

Given  $\Phi$ , we construct  $q_2$  as follows.

For each clause in  $\Phi$  we create an atom in  $q_2$ :

- If the clause is of Type 1, then the atom is  $R$ ;  
if the clause is of Type 2, then the atom is  $S$ ; etc.
- The variables in this atom are the same as the variables in the clause

---

Example:

$$\Phi = (\neg X_3 \vee \neg X_1 \vee X_4) \wedge (X_1 \vee X_2 \vee X_3) \wedge (\neg X_2 \vee \neg X_3 \vee X_1)$$

$$q_2 = S(x_3, x_1, x_4), K(x_1, x_2, x_3), S(x_2, x_3, x_1)$$

# Proof

## Claim:

There exists a homomorphism  $q_2 \rightarrow q_1$  iff  $\Phi$  is satisfiable

1. Suppose there is a satisfying assignment  $\theta$  for  $\Phi$ :
  - Define the function  $h: \text{Vars}(q_2) \rightarrow \text{Const}(q_1)$ :
    - If  $\theta(X_i) = 0$  then  $h(x_i) = 0$
    - If  $\theta(X_i) = 1$  then  $h(x_i) = 1$
  - $h$  is a homomorphism  $h : q_2 \rightarrow q_1$  (why??)
2. Suppose there exists a homomorphism  $h : q_2 \rightarrow q_1$ .
  - Define the following assignment  $\theta$ :
    - If  $h(x_i) = 0$  then  $\theta(X_i) = 0$
    - If  $h(x_i) = 1$  then  $\theta(X_i) = 1$
  - $\theta$  is a satisfying assignment for  $\Phi$  (why??)

# Proof

## Claim:

There exists a homomorphism  $q_2 \rightarrow q_1$  iff  $\Phi$  is satisfiable

1. Suppose there is a satisfying assignment  $\theta$  for  $\Phi$ :
  - Define the function  $h: \text{Vars}(q_2) \rightarrow \text{Const}(q_1)$ :
    - If  $\theta(X_i) = 0$  then  $h(x_i) = 0$
    - If  $\theta(X_i) = 1$  then  $h(x_i) = 1$
  - $h$  is a homomorphism  $h : q_2 \rightarrow q_1$  (why??)
2. Suppose there exists a homomorphism  $h : q_2 \rightarrow q_1$ .
  - Define the following assignment  $\theta$ :
    - If  $h(x_i) = 0$  then  $\theta(X_i) = 0$
    - If  $h(x_i) = 1$  then  $\theta(X_i) = 1$
  - $\theta$  is a satisfying assignment for  $\Phi$  (why??)

This completes the proof: query containment is NP-hard.

# Discussion

- We have a procedure to check if two queries are equivalent
- Will use to do query optimization (discussed next)
- NP-hard is not that scary today: SAT solvers can scale to tens of thousands of clauses

# Query Minimization

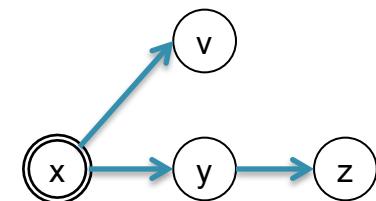
**Definition** A conjunctive query  $q$  is minimal if for every other conjunctive query  $q'$ , if  $q \equiv q'$  then  $q'$  has at least as many atoms as  $q$

# Query Minimization

**Definition** A conjunctive query  $q$  is minimal if for every other conjunctive query  $q'$ , if  $q \equiv q'$  then  $q'$  has at least as many atoms as  $q$

Are these queries minimal ?

$q(x) :- R(x,y), R(y,z), R(x,v)$

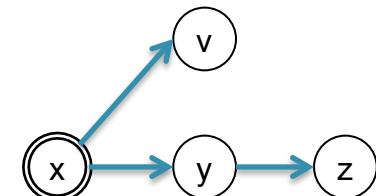


# Query Minimization

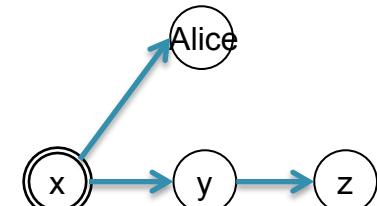
**Definition** A conjunctive query  $q$  is minimal if for every other conjunctive query  $q'$ , if  $q \equiv q'$  then  $q'$  has at least as many atoms as  $q$

Are these queries minimal ?

$q(x) :- R(x,y), R(y,z), R(x,v)$



$q(x) :- R(x,y), R(y,z), R(x, 'Alice')$



# Query Minimization

- Query minimization algorithm

Choose an atom of  $q$ , and remove it: let  $q'$  be the new query

We already know  $q \subseteq q'$  (why ?)

If  $q' \subseteq q$  then permanently remove it; otherwise try another atom

- Notice: the order in which we inspect atoms doesn't matter

# Query Minimization In Practice

- No database system today performs minimization
- Reason:
  - It applies only to SELECT-DISTINCT queries (doesn't work for bag semantics)
  - Users rarely write non-minimal queries
- However, non-minimal queries arise when using views intensively

# Query Minimization for Views

An Employee is a “happy boater” if she is a boater and her manager is also a boater:

```
CREATE VIEW HappyBoaters
```

```
SELECT DISTINCT E1.name, E1.manager  
FROM Employee E1, Employee E2  
WHERE E1.manager = E2.name  
    and E1.boater= ‘YES’  
    and E2.boater= ‘YES’
```

This query is minimal

# Query Minimization for Views

Now compute the Very-Happy-Boaters

```
SELECT DISTINCT H1.name  
FROM HappyBoaters H1, HappyBoaters H2  
WHERE H1.manager = H2.name
```

# Query Minimization for Views

Now compute the Very-Happy-Boaters

```
SELECT DISTINCT H1.name  
FROM HappyBoaters H1, HappyBoaters H2  
WHERE H1.manager = H2.name
```

After expanding HappyBoaters, the query is not minimal

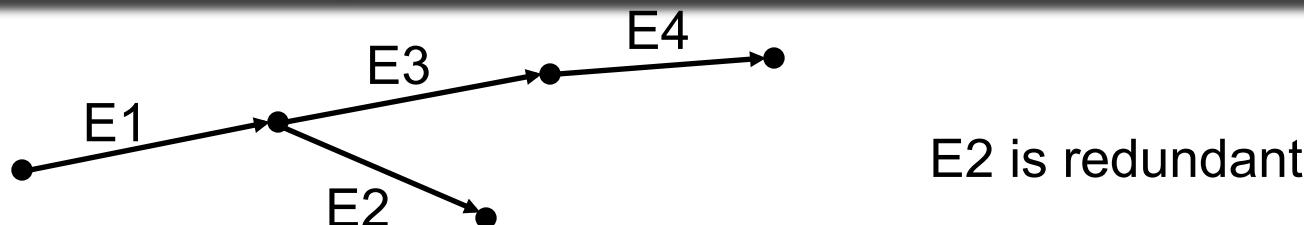
# Query Minimization for Views

Now compute the Very-Happy-Boaters

```
SELECT DISTINCT H1.name  
FROM HappyBoaters H1, HappyBoaters H2  
WHERE H1.manager = H2.name
```

After expanding HappyBoaters, the query is not minimal

```
SELECT DISTINCT E1.name  
FROM Employee E1, Employee E2, Employee E3, Employee E4  
WHERE E1.manager = E2.name and E1.boater = 'YES' and E2.boater = 'YES'  
and E3.manager = E4.name and E3.boater = 'YES' and E4.boater = 'YES'  
and E1.manager = E3.name
```



# Beyond CQ

- Any static analysis on relational queries (including containment, equivalence) is undecidable
- This follows from Trakthenbrot's theorem, discussed next
- Similar to Rice's undecidability theorem for Turing machines

# Trakhtenbrot's Theorem

**Definition** A sentence  $\varphi$ , is called finitely satisfiable if there exists a finite database instance  $D$  s.t.  $D \models \varphi$

**Theorem** [Trakhtenbrot] The following is **undecidable**: Given FO sentence  $\varphi$ , check if  $\varphi$  is finitely satisfiable

# Trakhtenbrot's Theorem

**Definition** A sentence  $\varphi$ , is called finitely satisfiable if there exists a finite database instance  $D$  s.t.  $D \models \varphi$

**Theorem** [Trakhtenbrot] The following is **undecidable**: Given FO sentence  $\varphi$ , check if  $\varphi$  is finitely satisfiable

Is this satisfiable?

$$\begin{aligned} & \exists x \exists y R(x,y) \wedge \\ & \forall y (\exists x R(x,y) \rightarrow \exists z R(y,z)) \wedge \\ & \forall x \forall y \forall z (R(x,y) \wedge R(x,z) \rightarrow y=z) \end{aligned}$$

# Trakhtenbrot's Theorem

**Definition** A sentence  $\varphi$ , is called finitely satisfiable if there exists a finite database instance D s.t.  $D \models \varphi$

**Theorem** [Trakhtenbrot] The following is **undecidable**:  
Given FO sentence  $\varphi$ , check if  $\varphi$  is finitely satisfiable

Is this satisfiable?

$$\begin{aligned} & \exists x \exists y R(x,y) \wedge \\ & \forall y (\exists x R(x,y) \rightarrow \exists z R(y,z)) \wedge \\ & \forall x \forall y \forall z (R(x,y) \wedge R(x,z) \rightarrow y=z) \end{aligned}$$

Is this satisfiable?

$$\begin{aligned} & \exists x \exists y R(x,y) \wedge \\ & \forall y (\exists x R(x,y) \rightarrow \exists z R(y,z)) \wedge \\ & \forall x \forall y \forall z (R(x,y) \wedge R(x,z) \rightarrow y=z) \\ & \exists x (\exists y R(x,y) \wedge \neg(\exists u R(u,x))) \end{aligned}$$

# Query Containment

**Theorem** Query containment for Relational Calculus is **undecidable**

**Proof:** By reduction from the finite satisfiability problem:

Given a sentence  $\varphi$ , define two queries:

$$q_1 = \exists x(R(x) \wedge \varphi), \text{ and } q_2 = \exists x(R(x) \wedge x \neq x)$$

Then  $q_1 \subseteq q_2$  iff  $\varphi$  is not finitely satisfiable

# Summary of DB Theory

- We have discussed query complexity and static analysis for conjunctive queries
- Why we care:
  - Complexity helps us understand the tradeoffs between what we can express in a language and how difficult it is to implement
  - Static analysis is critical in optimizing complex queries
- Next time: transactions